

○ USS TJR • PLATFORM AUDIT • 2026-08-22 • RERUN 2026-08-23

Command Deck Audit

Where Starship Endeavour's workbenches stand against the JARVIS-style personal command-centres and agentic ops platforms people are actually shipping in 2026 — an internal inventory pass, cross-checked against a real-world research pass, synthesized into named gaps.

Scope — 7 hub-listed workbenches, 6 unlisted, 18 non-UI platform capabilities

Method — internal file/code audit + external research, three independent passes

Update — 3 of 3 original recommendations shipped same day, live-verified against real data

Original verdict — substantial, real infrastructure underneath (Event Bus, Attention Engine, an 18-job scheduler, a local-first Model Router), but almost none of it is surfaced as a *command centre* in the sense the outside world now uses that word. Today, Endeavour is mostly a set of dashboards you visit — not a system that watches, decides, and acts on your behalf.

Same-day update — the two most load-bearing gaps are closed: scoring is real (Priority Engine) and the proactive layer has spoken for the first time (Attention Engine, real severity wiring). Notification duplication is consolidated. Health OSINT's real backlog (141 unreviewed signals) is cleared and now self-managing. A parallel investigation on Technical OSINT surfaced real misses (genuine data-breach stories getting suppressed) worth a look. Four gaps remain fully open: unified action layer, long-term memory, voice/ambient interface, agent observability.

Rerun, 2026-08-23 — same four gaps, sharper picture. Real benchmark numbers now exist for long-term memory (no system wins both accuracy and latency at once). MCP-shaped permission routing is the field's answer to a unified action layer, not a new platform to bolt on. And a fifth capability nobody named on the first pass turned up: a single "who's alive right now" panel over the Event Bus — the closest thing to a free win left on the table. Five build candidates below, ranked for a single delivery day.

CURRENT STATE

What's actually live

Seven workbenches reach the Captain through the hub tile grid. Five of those seven are read-only. Six more exist on disk but sit outside the hub entirely — reachable only by a direct URL, or not reachable at all.

WORKBENCH	MODE	PROACTIVE	LLM-SYNTHESIZED	STATUS
Captain's Chair	Read-only	No	Partial	live
Technical OSINT	Read-only	Fed by cron	Yes	live
Health OSINT	Read-only	Fed by cron	No	live
Content Workbench	Action-taking	No	Partial	live
Human Systems	Read-only	No	No	live
Advisory	Action-taking	No — by design	Yes	live
Briefs	Read-only archive	Fed by cron	Yes	live · redesigned this week

HUB-LISTED — reachable from /workbenches

WORKBENCH	MODE	NOTE
Captain's Brief Workbench	Read-only	Zero LLM calls, by explicit design — "every sentence traces to a countable fact"
Capture Workbench	Action-taking	Genuinely realtime — the one surface that reacts to something arriving, not a page load
Communications	Action-taking	direct-URL only — explicitly not in nav yet, per its own collection

	taking	rollout plan
Health OSINT Curation	Action-taking	Human-gated review queue, by design
Knowledge Workbench	Read-only	reduced scope — Library domain pulled back to draft mid-setup
Mission Workbench	Action-taking	Look-and-feel migration of legacy CRUD only

ON DISK, NOT HUB-LISTED – plus 27 legacy routes, most confirmed genuinely unreachable

The proactive layer exists — most of it has now actually spoken:

Attention Engine	Classifies events into interrupt_now / can_be_delayed / etc, running continuously.	fixed – real severity now drives real classification
Priority & Opportunity Engine	Feeds the "ranked" Decisions Inbox.	fixed – value_score went 0/200 → 95/200 events on tl
Event Bus	Cross-platform log, 12 real emit points, ~7,700+ rows.	real backbone
Unified Memory	One caller wired.	that caller is dead code – untouched
Notification service	Built, validated.	fixed – command_bus.py now sends through it, real c
Model Router	Local-first LLM gateway, cloud fallback.	reference-quality, per the platform's own registry
Scheduler	18 distinct jobs —	now 21 – daily digest, health auto-curation, suppr

	collection, briefs, QA, nudges, correlation.	
Search	Six separate implementations.	never consolidated – untouched
X0 Telegram bot	12 commands.	still pull/request-response – untouched

FIELD NOTES

What "command centre" means outside this platform

Two live research threads — consumer personal-AI-assistant products, and enterprise agentic-ops tooling — plus the failure modes both keep citing about themselves.

CONSUMER / PROSUMER

- **Limitless** (formerly Rewind's pendant) — near-perfect passive recall, explicitly does not predict or suggest. Acquired by Meta, device pulled from sale, app sunsetting.
- **Friend** — deliberately rejects productivity framing entirely; companionship only, no task execution.
- **Catch** — the named differentiator is "acts without prompting": watches calendar/inbox, initiates, lives in iMessage/text.
- **Home Assistant Assist** (2026) — fully local voice pipeline, real device actions in 2–4s.
- **Mem0 / Letta / Zep / LangMem** — long-term personal memory is now its own infrastructure category: three distinct layers, not one table.

ENTERPRISE / AGENTIC OPS

- **Microsoft Security Copilot** — autonomously triages alerts, offers one-click containment across the SOC lifecycle. Closed-loop, not just flagged.
- **Sierra** — explicitly "agents with memory, context, and the ability to act across systems," contrasted against knowledge-only platforms like Glean.
- **Agent observability** (Braintrust, LangSmith, Arize, Datadog LLM Observability) — a named 2026 category, because a single multi-agent request can produce 40–200 spans no human can read raw.

WHAT THEY SAY FAILS

- **Capture ≠ proactivity** — Limitless is the cited case: great memory, zero anticipation, discontinued.
- **Alert fatigue** is named as the central design problem in both consumer and SOC contexts — severity tiering and dedup are the standard fix.
- **"Talks about your life vs. lives your life"** — always-listening capture still needs the human to ask the right question.

GAP ANALYSIS

Seven named gaps – two resolved, one partial

Not ranked — each stands on its own evidence, from the audit on one side and the field notes on the other. Resolved/partial gaps keep their original framing, with a same-day resolution block underneath.

No unified action layer

structural – still open

Five of the seven hub-listed workbenches are pure dashboards. Only Content, Capture, and Mission actually change state on your behalf.

HERE

Captain's Chair, both OSINT workbenches, Human Systems, and Captain's Brief Workbench are all request-triggered and read-only.

ELSEWHERE

Sierra and Security Copilot define themselves by the ability to act across systems — explicitly the line that separates them from knowledge platforms like Glean.

Precedent extended, not the platform gap itself Health OSINT's real backlog (141 unreviewed signals, one Sunday-night review step that never once actually happened) is now auto-decided by the same publish/reject action pattern Content Workbench already proved — 118 published, 18 rejected, live.

That's one more workbench with a real action loop. not the unified layer this gap actually names.

The proactive layer is invisible and silent

resolved 2026-08-22

The infrastructure to notice things on its own exists. It has never actually spoken up.

HERE (ORIGINAL)

Attention Engine's interrupt_now bucket has fired zero times on real data. The only unprompted contact with the Captain is four scheduled Telegram sends a day.

ELSEWHERE

"Acts without prompting" is Catch's entire pitch. Limitless — recall without anticipation — was acquired and discontinued.

Resolution Root cause traced past the threshold (which was correct) to the input: every event's importance was rank_score, deliberately deflated for brief-ranking, never observed above 70.80 across 1,000 real events against a 75 floor. Replaced with a 4-tier severity score from fields the classifier already computes. A drill mechanism (weekly, real Telegram send) proves the delivery pipe end to end; real historical events (a Google Cloud Status incident, a DocuSign incident, an APRA Banking Act update) now correctly cross the threshold when run through the real engine. A first-draft formula that over-fired (~1 every 2.5 hours) was tested and rejected in favor of one that fires ~1 every 4.3 days — deliberately tuned against the exact alert-fatigue failure this audit named.

No real long-term memory infrastructure

architectural

Data lives in per-domain tables. Nothing resembles a structured, cross-session context graph.

HERE

Unified Memory has exactly one caller wired up, and that caller is itself dead code.

ELSEWHERE

Mem0, Letta, Zep, and LangMem are named as the entire 2026 category for this — three distinct memory layers, purpose-built, not a side effect of a database schema.

Notification design has no real discipline

partially resolved 2026-08-22

The pieces that should decide what deserves your attention were built once and then bypassed

HERE (ORIGINAL)

Notification service: built, validated, zero production callers. `command_bus.py` sends ad hoc instead.
Search: six separate implementations, never consolidated.

ELSEWHERE

Severity tiering and dedup are named as the standard fix to the single most-cited failure mode across both consumer and SOC tooling in 2026.

Resolution — notification duplication only `command_bus.py`'s real ALERT/CRITICAL sends (all 3 live call sites) now go through `notification_service.notify()` instead of its own raw senders — one implementation instead of two, with real retry/backoff/call-log. Building this also exposed and fixed a genuine, separate production bug: legacy Telegram Markdown was silently failing to render escaped underscores as literal characters instead of consuming them — caught only because a live send was checked on a real device, not trusted from an "ok: true" API response. Fixed by switching to HTML `parse_mode` platform-wide. Search's six implementations remain unconsolidated — out of scope for a quick win.

No voice or ambient interface

interface

Every surface on this platform is pull — open a page, or read a bot reply. Nothing listens.

HERE

Zero voice/ambient input anywhere in the stack.

ELSEWHERE

Home Assistant's fully local 2026 voice pipeline (Wyoming + whisper + Piper + Ollama) does real device actions in 2–4 seconds — proof this is achievable at hobbyist scale, not just by well-funded labs.

Multi-agent work happens with zero observability

operational

There's no trace UI anywhere. When something goes wrong across agents, the only tool is raw logs.

HERE

This week's 8-agent Mistral pipeline work traced a hallucination bug by hand, through raw `systemd` journal output — no dashboard exists for "what did the AI actually do and why."

ELSEWHERE

Agent observability (Braintrust, LangSmith, Arize, Datadog LLM Observability) is now a named enterprise category, specifically because raw agent logs stop being legible past a handful of hops.

Scoring is fake where it matters most

resolved 2026-08-22

Not a comparison gap — an internal one, but it undercuts every other claim of "the system decided this matters."

HERE (ORIGINAL)

Priority & Opportunity Engine's value / opportunity / time_sensitivity dimensions are hardcoded to 0. The "ranked" Decisions Inbox isn't actually ranked by anything real yet.

WHY IT COMPOUNDED

Every gap above assumed the system can tell what matters. This was the piece that made that assumption false.

Resolution Root cause was wiring, not the engine — score_event() already computed correctly when given real inputs. Its one real caller was simply never passing value_dimensions or opportunity_value at all. Fixed by deriving a first-pass value signal from each event's existing categorical fields, and wiring a real time_sensitivity urgency curve into the one emit point with a genuine due date (missions). Live-verified against real core_events: value_score > 0 went from 0/200 events to 95/200 on the same call. opportunity_score correctly stayed at 0/200 — no content-intelligence events in the current sample, honestly absent rather than fabricated.

FAIR COMPARISON

Where this is already ahead

Not everything is a gap. A few things here are genuinely rarer than what the named examples above are shipping.

Cross-domain breadth under one roof. Health, engineering, intelligence, and mission tracking share one real Event Bus — 12 emit points, ~7,700+ rows. Every named consumer product above is single-purpose: Limitless does memory, Friend does companionship, nothing else.

A genuinely local-first LLM gateway. The Model Router (cloud fallback, not cloud-first) is called reference-quality by the platform's own registry — most hobbyist "build your own JARVIS" projects are still working this part out.

A real closed-loop content pipeline. This week's daily-digest redesign — auto-publish, multi-agent narrative synthesis, a grounding safeguard traced and fixed against a live hallucination bug — is a rarer thing than it looks, even against the named examples here.

A second real closed-loop decision system, same day. Health OSINT's auto-curation makes the same PUBLISH/REJECT/ESCALATE call Content Workbench proved, but adds a real safety valve: genuinely ambiguous cases are left for a human rather than forced either direction — "default to ESCALATE over guessing, this is health data" wasn't a suggestion, it changed the actual decision logic.

Willing to find its own real misses. A read-only audit built for Technical OSINT immediately surfaced genuine false suppressions — real data-breach stories the mechanical filter dropped. Most systems at this stage would rather not look.

SHIPPED

All 3 original recommendations, same day

What was recommended, and what actually happened when it was built and live-verified.

- ✓ **Fixed the scoring.** value_score went 0/200 → 95/200 events on the same live call. The wiring gap was a caller never passing two fields, not a broken engine.
- ✓ **Gave the Attention Engine its chance to speak up.** A drill mechanism proved real Telegram delivery end to end — and in doing so, found and fixed a second real bug (Markdown escaping silently rendering broken) that would have swallowed a genuine alert even after everything else was fixed. Real severity now wired into every classified event; real historical events correctly cross the interrupt threshold.
- ✓ **Made a workbench act, not just show.** Health OSINT's 141-signal backlog — unreviewed since the feature launched — is now auto-decided (118 published, 18 rejected) with a real escalation path for anything genuinely ambiguous, and stays clear going forward via the existing weekly job.

Also found along the way, not on the original list: a read-only suppression audit for Technical OSINT surfaced real data-breach stories getting mechanically filtered out. Two independent Gemini bugs (a rejected request field, a wrong credential) — neither was actually the billing issue it looked like. And the clearest lesson of the day: an API returning `ok: true` proves a message was accepted, never that it rendered correctly — caught only because a real device was checked twice, not trusted from the API response alone.

Four gaps remain fully open: no unified action layer, no long-term memory infrastructure, no voice or ambient interface, and multi-agent work still has zero observability. None of them are quick wins — each is a real, multi-day build.

RERUN • 2026-08-23

What moved in the field since the first pass

Same four open gaps, but the outside world got more specific about how to close them. Three findings worth knowing before picking a dependency.

MEMORY HAS REAL NUMBERS NOW

LongMemEval benchmark, 2026 — no system leads on both axes:

SYSTEM
Hindsight + Gemini
Supermemory
Zep + GPT-4o
Mem0 + GPT-4o

ACTION LAYER HAS A NAME NOW

- Tiered autonomy** is the 2026 standard framing — 3 tiers sized to recovery cost if the action is wrong, not to what the action is. Tier 1 runs itself; Tier 2 asks async; Tier 3 needs a sync sign-off.
- MCP as action router** — a permission check sitting between agent and tool call is now the dominant pattern, not a bespoke platform. A bridge, not a rebuild.
- "Human as signal, not gate"** — 2026 practice rejects blocking yes/no approval; the agent reasons over pending-approval state and resumes when it clears.

AN UNNAMED GAP SURFACED

- Hobbyist "agentic OS command centre" writeups converge on four non-negotiable panels: **every agent alive right now**, cross-agent memory search, a goals/journal log, and a cost/token analytics view.

THREE NAMED REPOS, CROSS-CHECKED

- LifeOS** — its **Pulse** daemon runs voice, hooks, observability, cron, and dashboards as one process, not five. Its **Cortex** memory is exactly the shape item 04 below proposes: a hot-layer plus a typed archive, not a vector store alone.